

BEAST Lab Organization

LRZ

Dr. Josef Weidendorfer, josef.weidendorfer@lrz.de

Amir Raoofy, amir.raoofy@tum.de

LMU


Sergej Breiter, sergej.breiter@nm.ifi.lmu.de

Minh Thanh Chung, minh.thanh.chung@ifi.lmu.de

Dr. Karl Fűrlinger, karl.fuerlinger@ifi.lmu.de

TUM

Bengisu Elis, bengisu.elis@tum.de



TUM Uhrenturm



Table of Contents

Course Organization

Introduction to BEAST

OpenMP Introduction

Tentative Course Structure

- 12 meetings in total
 - Last meeting on January 25th 2024
- 7 Assignments (Assignment 0 is an informational handout)
 - 1 week each
- 2 Projects
 - 2 weeks each
- Student groups of 3 BA students and 2 MA students
- Two groups will present their reports at the meeting right after deadline
 - Presentation notification will be sent 2 days before presentations
 - No slides, go through your report and talk about your findings
 - About 15-20 minutes talk per group
 - Focus on the **most important and interesting** aspects of your report
 - You **don't** need to address every part of your report
 - Q/A after the student presentations (everyone is highly encouraged to ask questions)

Repository Structure

Gitlab main repository: <https://gitlab.lrz.de/beastlab23ws>

- Each team has a repository, which includes:
 - Lecture slides
 - Assignment material
 - Code template
 - Your code submissions and report (once you place it there)
- Only solutions on the `main` branch will be graded !
 - At the due date, your current master branch state is automatically tagged and archived
 - Make sure your code and report is there by the deadline
- Machine account information will be sent via e-mail - Change passwords please !
- We need you to sign your commits.
- Tutorial: https://docs.gitlab.com/ee/user/project/repository/gpg_signed_commits/
- If GitLab shows a **Verified** label on your commits you are good to go.
- We will tag your last commit before each deadline.

Infrastructure Usage

Two-Factor Authentication

- You will work on LRZ systems as part of the lab.
- LRZ requires 2FA.
- Today, We will make sure that everybody registers a 2nd factor.
- For this, we expect you to have a smart phone with an authenticator app.
 - privacyIDEA authenticator
 - or Microsoft Authenticator
- Instructions: <https://doku.lrz.de/two-factor-authentication-prerequisites-35882365.html>

Access to BEAST and each individual system

- see Assignment 0

Infrastructure Usage

Exclusive Resource Allocation

- We use SLURM this semester to schedule your jobs and measurements.
- For your development throughout the week, you can directly ssh into individual machines
- For final measurements we allocate 1 full day where no direct ssh can be done.
- **this semester we allocate Tuesdays of every week for exclusive slurm jobs**
- In this timeframe you can only run code through slurm allocations
- We limit the job timelimit to 15 minutes to ensure every group gets a chance to schedule jobs during that day.
- here are some usefull links:
`https://doku.lrz.de/example-parallel-job-scripts-on-the-linux-cluster-10746636.html`
- usefull commands: `sinfo`, `squeue`, `salloc`, `srun`, and `sbatch`

Up Next: Introduction to BEAST

BEAST Lab

Introduction to Systems

October 19, 2023 | Josef Weidendorfer

Collaboration among 3 institutions

LMU
TUM
LRZ

LMU – MNM/Prof. Kranzlmüller
(Karl Förlinger, Minh Chung, Sergej Breiter)

TUM – CAPS/Prof. Schulz
(Bengisu Elis)

LRZ - Future Computing Group
(Josef Weidendorfer, Amir Raoofy)

We want you to learn about **performance properties of current architectures**

- Be able to understand and explain performance effects seen from measurements
- Get a deeper understanding of current system designs (CPU / GPU)

Part 1: get started with small codes across systems

- We show key hardware design concepts + a parallel programming model (OpenMP)
- We give you typical small HPC / microbenchmark code examples
- You run measurements of different scenarios across systems, compare / discuss results
- We all discuss results in weekly meetings, starting with presentations of groups

Structure:

CPU evaluation (Memory, Compute) → GPU evaluation → Tool

We want you to learn about **performance properties of current architectures**

- Be able to understand and explain performance effects seen from measurements
- Get a deeper understanding of current system designs (CPU / GPU)

Part 2: make use of gained knowledge – Final Project

- We assign randomly one system to each group
- We give you some larger typical HPC code
- You tune the code to get best single-node performance (3 week time)
- We discuss intermediate/final experiences/results in weekly meetings


Target Architectures for the Lab

CPUs

- Intel Icelake (ISA: x86-64 + AVX512)
- AMD MilanX (ISA: x86-64 + AVX2)
- Marvell ThunderX2 (ISA: ARM AArch64 + Neon)
- Fujitsu A64FX (ISA: ARM AArch64 + SVE)

GPUs

- NVidia V100 (eventually also A100)
- AMD MI-100 (eventually also MI-210)



SuperMUC-NG
Top500 (Nov 2018): #8
Lenovo Intel (2019)
311,040 cores
Intel Xeon Skylake
26.9 PetaFlops Peak
19.5 PetaFlops Linpack
719 TeraByte Main Memory
70 PetaByte Disk

BEAST – Bavarian Energy Architecture and Software Testbed

The LRZ Future Computing Testbed



- Help decide about next large system
 - Get experience on benefits of various future architectures for LRZ codes
 - Find best configuration: how much money to spend on compute / memory / network?
 - Enable migration planning: educate own staff / port LRZ tools / prepare courses
 - Support vendor collaboration
- Enable research studies on new technologies
 - Forward looking: LRZ services around future platforms, novel usage models
 - more experimental: FPGAs, AI accelerators, integration of heterogeneity (QC)
 - In partnership with selected researchers from Munich universities

Lot of work to do! Engage students for student work (BA, MA): This Lab!

The Testbed – Available Hardware



2 racks, each with 6 PDUs (for power measurements)

- Max power consumption per rack: 35 kW

Top to bottom (picture from last year)

- 3 switches (Infiniband 200Gb/s HDR), 2x 48port 1Gb/s Ethernet
- Login 1U “testbed.cos.lrz.de”
- **2x AMD Rome** GPU server 2U: “rome1” / “**rome2**”
(to be upgraded: AMD MilanX + GPU)
- Storage 2U with homes
- **2x Marvell ThunderX2** GPU server 2U: “thx1” / “**thx2**”

Not shown:

- HPC CS500 Management 2U + **8 nodes A64FX** “cs1” – “cs8”
- **2x Intel IceLake** GPU server 2U: “ice1” / “**ice2**”



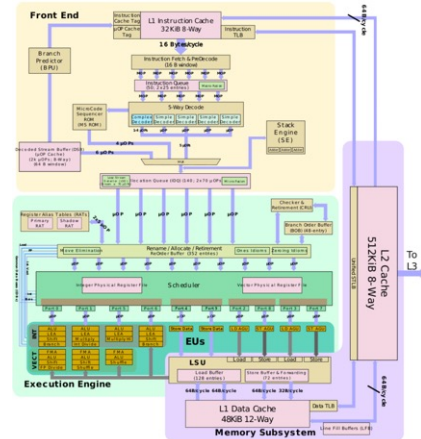
Intel Icelake

Two systems in BEAST

- 2 sockets Intel Xeon (Icelake) Platinum 8360Y
 - 2x 36 = 72 cores
 - 2x 512bit vector units per core (8 x DP FMA)
 - 2 threads per core ("Hyper-Threading")
 - 2.4 GHz base, Intel 10nm
- 512 GB main memory, 1.5 TB Optane NVRam

Links

- [https://en.wikichip.org/wiki/intel/microarchitectures/ice_lake_\(server\)](https://en.wikichip.org/wiki/intel/microarchitectures/ice_lake_(server))
- https://en.wikichip.org/wiki/intel/microarchitectures/sunny_cove

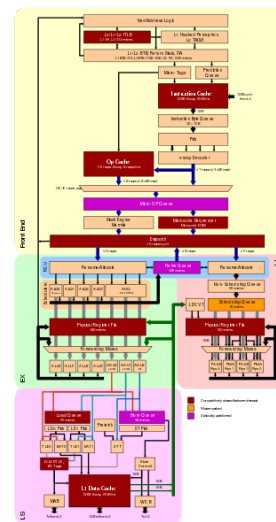



Two systems in BEAST

- 2 sockets with EPYC 7773X
- 2x 64 = 128 cores (“Zen3”)
 - Chiplet design: IO-Die + 8x CCX-Dies (each 8-core)
 - 2x 256-bit vector units per core (4 x DP FMA)
 - 2 threads per core
 - 2.2 GHz base, TSMC 7nm
- 1 TB main memory
- 8x AMD Radeon MI-210 GPUs
 - 7nm, 64GB HBM, PCIe4

Link

- https://en.wikichip.org/wiki/amd/microarchitectures/zen_3



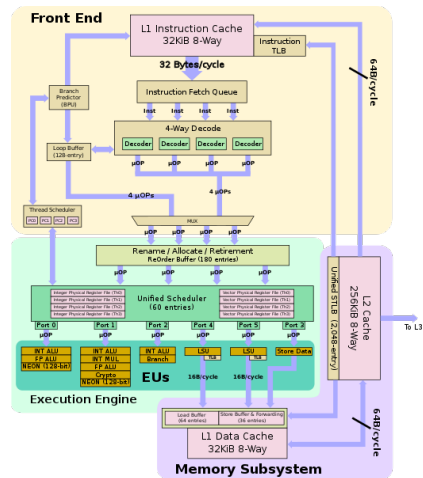
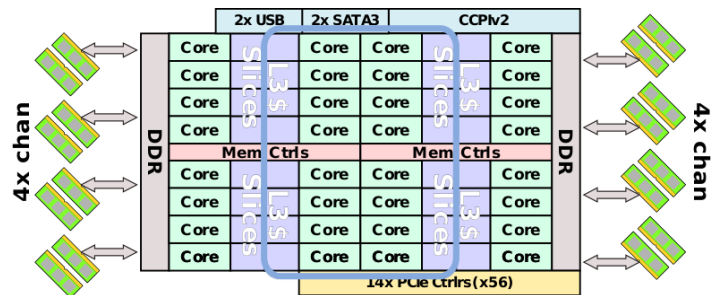
Marvell ThunderX2

Two systems in BEAST

- 2 sockets with ThunderX2 CN9980
- 2x 32 = 64 cores (“Vulcan”)
 - 128-bit vector units (2 x DP FMA)
 - 4 threads per core
 - 2.2 GHz base, 16nm
- 512 GB main memory
- 2x Nvidia V-100
 - Volta, 32GB HBM, PCIe3

Link

- <https://en.wikichip.org/wiki/cavium/microarchitectures/vulcan>




Fujitsu A64FX

HPE CS500 in BEAST

- 8 nodes with one A64FX CPU (“NSP1”)
- 48 cores per CPU
 - 2x 512bit vector units per core
 - 1.8 GHz, TSMC 7nm
 - 4 NUMA domains
- 32 GB HBM2

Link

- https://en.wikipedia.org/wiki/Fujitsu_A64FX



[Fujitsu: The 1st SVE Enabled Arm Processor: A64FX and Building up ARM HPC Ecosystem, 2019]

Access via Linux Cluster login nodes

- `ssh XXX@lxlogin1.lrz.de` (or `lxlogin2` / `3` / `4`)
 - need to configure 2FA TOTP/PUSH via <https://simmfa.sim.lrz.de>
- `ssh testbed.cos.lrz.de`
- `ssh <system>`

If `testbed.cos.lrz.de` is not reachable, retry after 1 hour

- probably just a reboot

Compilers

- system: “gcc”
- via modules: see “module avail”, then “module load <package>”

Access and Usage: Intel Icelake @ BEAST



Access

- `ssh XXX@lxlogin1@lrz.de`
- `ssh testbed.cos.lrz.de`
- `ssh ice1`

Compilers

- `gcc, icc` (Intel compiler)

Access and Usage: AMD Milan-X @ BEAST



Access

- `ssh XXX@lxlogin1.lrz.de`
- `ssh testbed.cos.lrz.de`
- `ssh milan2`

Compilers

- `gcc`, `clang` (from AMD RocM)

Access and Usage: ThunderX2 @ BEAST



Access

- `ssh XXX@lxlogin1.lrz.de`
- `ssh testbed.cos.lrz.de`
- `ssh thx2`

Compilers

- `gcc`
- (via „module load cuda/11.1.1 llvm“) `clang`

Access

- `ssh XXX@lxlogin1.lrz.de`
- `ssh testbed.cos.lrz.de`
- `ssh cs1 / cs2`

Compilers

- `gcc (8)`
- `gcc 11` and `13` (via „`module load gcc/11.0.0` or `module load gcc-13.1.0-djjcsa ...`“)
- Cray compiler: “`cc`”, enable OpenMP: “`-h omp`”



Leibniz Supercomputing Centre
of the Bavarian Academy of Sciences and Humanities

Up Next: OpenMP Introduction

Sergej Breiter (presenting) and Dr. Karl Furlinger

Institut für Informatik

Lehrstuhl für Kommunikationssysteme und
Systemprogrammierung

OpenMP Basics

BEAST Lab WS 2023/24

Praktikum

Evaluierung moderner HPC-Architekturen
und -Beschleuniger



- A method for portable programming of shared memory systems
 - **Open** specification for **Multi-Processing**
- Industry Standard
 - Guided by the OpenMP **Architecture Review Board** (ARB)
 - Major companies and research labs participate in the ARB
 - Current version: v5.2 (November 2021)
- Language extension for C/C++ and Fortran
 - Compiler **directives**
 - Library **routines**
 - Environment **variables**
- www.openmp.org
 - Current specification, tutorials, other resources (such as examples)



OpenMP Example: Hello World

Source Code:

```
#include <stdio.h>
#include <omp.h>

int main(int argc, char* argv[]) {
    #pragma omp parallel
    {
        printf(„Ahoi OpenMP world\n“);
    }
}
```

Compilation:

```
icc -qopenmp hello.c -o hello
```

```
icx -fopenmp hello.c -o hello
```

```
gcc -fopenmp hello.c
```

The flag to enable OpenMP is
implementation-specific


Execution:

```
>$ export OMP_NUM_THREADS=4
>$ ./hello
Ahoi OpenMP world
Ahoi OpenMP world
Ahoi OpenMP world
Ahoi OpenMP world
```

Execution with 2 threads:

```
>$ export OMP_NUM_THREADS=2
>$ ./hello
Ahoi OpenMP world
Ahoi OpenMP world
```

OpenMP Execution Model

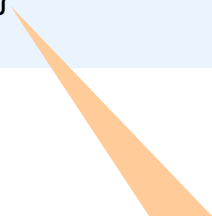


- This model is called the **fork-join Model**
 - Program starts with a single thread (called the **initial thread**)
 - Parallel regions create additional threads (**team threads**), initial thread becomes the **master thread** in the team
 - Team threads disappear (logically) at the end of a parallel region
 - Implementations may keep team threads around in a thread pool for reasons of efficiency
 - There is an **implicit barrier** at the end of a parallel region
 - Number of threads **may change** between parallel regions

Creating Threads: #pragma omp parallel

```
#include <stdio.h>
#include <omp.h>

int main(int argc, char* argv[]) {
    #pragma omp parallel
    {
        printf(„Ahoi OpenMP world\n“);
    }
}
```



The structured block is
executed redundantly
(in parallel) by all
threads

■ **Worksharing** constructs are
used to distribute work
between threads

- for
- sections
- single
- workshare (Fortran only)

Shared and Private Variables

- Variables declared outside the parallel region are **shared** by default

```
#include <stdio.h>
```

```
#include <omp.h>
```

shared by default

```
double alpha=1.23;
```

```
int main(int argc, char* argv[]) {
```

```
    double gamma=23.11;
```

shared by default

```
    #pragma omp parallel
```

```
    {
```

```
        int mydelta;
```

private by default

```
        #pragma omp for
```

```
            for(int i=0; i<100; i++) {
```

```
                do_some_work(i, alpha);
```

```
            }
```

```
        mydelta = ...;
```

OK! modifying private copy

```
        gamma+=mydelta;
```

```
    }
```

```
}
```

!!!Warning: modifying shared variable!!!
Needs some form of synchronization, e.g.,
atomic, critical, locks

- Shared variables

- Are accessible by all threads
(only one copy exists)

- Private variables

- Accessible only by one
thread (each thread has its
own copy)

- Data sharing clauses can
override defaults

Data Sharing Clauses (Parallel and Work Sharing Constructs)

- **private**(var-list)
 - Variables in var-list are **private**
- **shared**(var-list)
 - Variables in var-list are **shared**
- **default**(private | shared | none)
 - Sets the default for all variables in this region
 - Default **none** raises compiler error if sharing is not explicitly specified
- **firstprivate**(var-list)
 - Variables are private and are initialized with the value of the shared copy before the region
- **lastprivate**(var-list)
 - Variables are private and the value of the thread executing the last iteration of a parallel loop in sequential order is copied to the variable outside of the region.

Initialization of Private Variables

```
int i, j;  
i = 1;  
j = 2;  
  
#pragma omp parallel private(i) firstprivate(j)  
{  
    printf(„i=%d j=%d\n“, i, j);  
}
```

Execution:

```
>$ export OMP_NUM_THREADS=4  
>$ ./a.out  
i=5456498 j=2  
i=-732837541 j=2  
i=788564 j=2  
i=821656 j=2
```

- Private copies of i are **not initialized!**
- Firstprivate copies of j are **initialized to the outside value**

Worksharing in Parallel Regions

- Goal: distribute work among threads in a parallel region

```
#include <stdio.h>
#include <omp.h>

int main(int argc, char* argv[]) {
    int i;
    #pragma omp parallel
    {
        #pragma omp for
        for(i=0; i<100; i++) {
            do_some_work(i);
        }
    }
}
```

```
// shorthand notation for the above
// combined parallel-workshare
#pragma omp parallel for
```

■ The `omp for` construct

- Specifies that the work (loop iterations) should be distributed to the available threads
- **Asserts** that the loop iterations are independent and can be parallelized

Parallel Loop (C/C++)

```
#pragma omp for [clause[,] clause]  
for(i=0; i<..; i++..) { .. }
```

- Loop iterations are distributed between the threads of the team
 - A **loop scheduling clause** specifies exactly **how**
 - Loop scheduling options: **static**, **dynamic**, **guided**, **auto**, and **runtime**
 - E.g., *schedule(static)*
- Characteristics:
 - There is **no synchronization (i.e., barrier) at the entry** of the loop
 - There is an **implicit barrier** at the end of the loop unless a *nowait* clause is specified
 - The loop iteration variable is **private by default**
 - Only **simple** (so-called **canonical forms**) of loops are supported
 - Integer iteration variable, only modified in the increment expression
 - Iteration count can be computed before executing the loop


```
#pragma omp for schedule(type[, size])
```

- Scheduling type is one of:
 - **static**: chunks of iterations of the specified size are distributed among threads in a **round-robin** fashion
 - **dynamic**: Threads **request chunks** of the specified size from the runtime; when finished executing, a thread requests a new chunk
 - **guided**: like dynamic, but the chunk size is proportional to remaining work; size parameter specifies the minimal chunk size
 - **auto**: decision is delegated to the compiler and/or runtime system
 - **runtime**: defer scheduling decision to runtime selection (via environment variable **OMP_SCHEDULE**); note that it is only possible to specify one schedule for all loops via an environment variable

Loop Scheduling Example

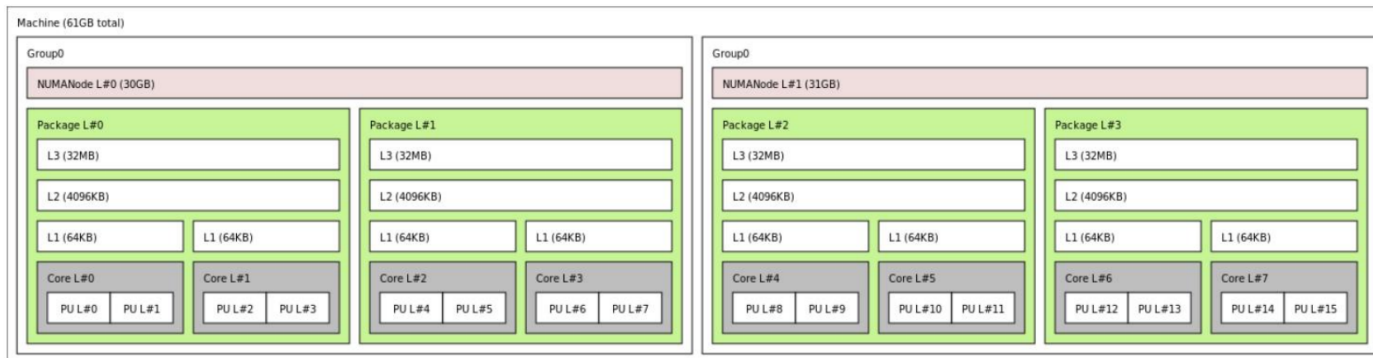
```
#pragma omp parallel
{
  //#pragma omp for schedule(static)
  //#pragma omp for schedule(static,3)
  //#pragma omp for schedule(dynamic,1)
  #pragma omp for schedule(dynamic,3)
  for(i=0; i<20; i++) {
    do_some_work(i);
  }
}
```

 Thread 0
 Thread 1
 Thread 2
 Thread 3



Thread Affinity

- How threads are mapped to hardware may influence performance
 - E.g., placement of threads to optimize cache sharing vs. memory bandwidth
 - HWLoc output example



- OpenMP allows the specification of
 - What we consider the unit of locality
OMP_PLACES env. variable = **threads** | **cores** | **sockets**
 - How to distribute threads to places
 - **OMP_PROC_BIND** env. variable and **proc_bind** clause
master | **spread** | **close**

- **OMP_PLACES** specifies a list of places where threads should be executed
 - **sockets** – each place corresponds to a single socket, a socket can have multiple cores
 - **cores** – each place corresponds to a single core, each core can have multiple hardware threads
 - **threads** – each place corresponds to a single hardware thread

export OMP_PLACES=cores

- Places and place lists can also be specified numerically
 - Meaning of numeric IDs depends on the system (/proc/cpuinfo, lscpu)

export OMP_PLACES={0,1,2,3}, {4,5,6,7}, ...

- **OMP_PROC_BIND(policy)** or **proc_bind(policy)** clause specify how threads are mapped onto places
 - **master** – each thread in the team is assigned to the same place as the master thread
 - **close** – threads in the team are placed close to the master thread
 - **spread** – threads are spread evenly over the places

- **Examples (HW as in Hwloc example)**

Parallel region with two threads, one per socket

OMP_PLACES=sockets

`#pragma omp parallel num_threads(2) proc_bind(spread)`

Parallel region with four threads, all on one socket

OMP_PLACES=cores

`#pragma omp parallel num_threads(4) proc_bind(close)`

Optimizing for NUMA (1)

- NUMA=Non-Uniform-Memory Access
 - Accessing local data is beneficial for performance
 - Virtual memory is mapped to physical memory in the granularity of pages (typically 4KB)
 - Usually where a memory page gets allocated is determined by a **first touch policy** (i.e., local to the core that first uses a memory page)
 - This implies that the initialization of data structures should reflect the intended later access patterns
 - Bad: Serial initialization and parallel access
 - Bad: Different parallel initialization and parallel access
 - Good: Parallel initialization and parallel access in same way
- Other options:
 - Explicit control using OS mechanisms, e.g., **numactl**

Optimizing for NUMA (2)

- Bad: serialized initialization leads to allocation of B,C,D all in one locality domain


```
// initialize data structures
for(i=0; i<N; i++ ) {
    B[i]= . . .
    C[i]= . . .
    D[i]= . . .
}
```

```
#pragma omp parallel for
for( i=0; i<N; i++ ) {
    A[i] = B[i]+C[i]*D[i];
}
```

- Good: parallel initialization in the same way it is later accessed (distributed across locality domains)

```
// initialize data structures in parallel
#pragma omp parallel for
for(i=0; i<N; i++ ) {
    B[i]= . . .
    C[i]= . . .
    D[i]= . . .
}
```

```
#pragma omp parallel for
for( i=0; i<N; i++ ) {
    A[i] = B[i]+C[i]*D[i];
}
```



ccNUMA system with
four locality domains

Image source: Hager, Wellein:
"Introduction to High Performance
Computing for Scientists and
Engineers"